A decorative graphic on the left side of the slide consisting of two overlapping parallelograms. The front one is blue and the back one is a light green color. They are positioned diagonally, with the blue one in front of the green one.

Learning to play Go

Vsevolod Karpov & Nicolas B. Carbone

The game of Go

- Go is a board game. Very Popular in Asia
- Played on 19x19 (most popular) board.
- The objective of the game:
 - Capture most territory
- After 100 moves:
 - ca. $1 \cdot 10^{(249)}$ possible positions
 - Chess: ca. $1.6 \cdot 10^{(160)}$
- A heuristic approach will not work.
Too many possible moves!

Main Goal:

- We want to train a network to master a 9x9 board.



Previous work - AlphaGo Zero

- Improved version of AlphaGo
- NN that beats best human players
- Pure reinforcement learning is better than learning from human games (AlphaGo), to achieve superhuman levels
- Uses Monte-Carlo-Tree-Search to explore the best predicted moves
- Also has a raw network which does not use any form of tree search - performs fairly well



AlphaGo vs Lee Sedol in
Seoul march 2016



Reinforcement Learning -Our approach

- 1) Start N networks with random weights.
- 2) Let each network play 2 matches against every other network (one as black and one as white).
- 3) Choose best N/2 networks based on win ratio (reward function).
- 4) Delete worst N/2. Clone best N/2.
- 5) Adjust weights of all networks by a random number in range $(-1,1) \cdot \text{learning_rate}$
- 6) Repeat from 2)
 - Current board and previous 6 boards are used as input
 - No tree search for future moves.



Network architecture : Stage 1

- Conv2d (features=32, stride=1, kernel=3, padding=1) , ReLU , BatchNorm
- Conv2d (features=64, stride=1, kernel=3, padding=1) , ReLU , BatchNorm
- Conv2d (features=128, stride=1, kernel=3, padding=1) , ReLU
- MaxPool(kernel=2, stride=2, padding=0) , BatchNorm
- Conv2d (features=256, stride=1, kernel=3, padding=1) , ReLU
- MaxPool(kernel=2, stride=2, padding=0) , BatchNorm
- Linear(82)
- Linear(82)
- Linear(82)
- Linear(82)



Network : Stage 1

- Stage 1 was meant to get a feeling of training time and progress
- Learning rate : 0.005, $N = 10$
- Trained for 500 epochs.
 - Achieved 100% win ratio against random play
- Used for validation for stage 2.




Network architecture : Stage 2

(Similar to AlphaGo)

- Sequential layer x7
 - Conv2d (features=64, stride=1, kernel=3, padding=1) , ReLU , BatchNorm, x2
 - Passthrough
- Conv2d (features=128, stride=1, kernel=3, padding=1) , ReLU , BatchNorm
- Conv2d (features=256, stride=1, kernel=3, padding=1) , ReLU
- MaxPool(kernel=2, stride=2, padding=0) , BatchNorm
- Conv2d (features=512, stride=1, kernel=3, padding=1) , ReLU
- MaxPool(kernel=2, stride=2, padding=0) , BatchNorm
- Linear(656)
- Linear(328)
- Linear(164)
- Linear(82)
- Softmax


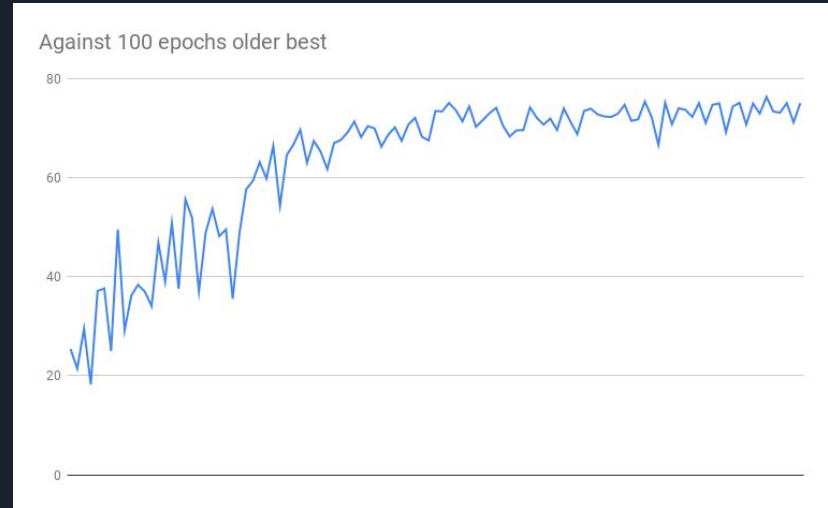
Stage 1 vs Stage 2 validation

- Blue :
 - learning_rate = 0.0005
 - N = 10
- Red :
 - Learning_rate = 0.0001
 - N = 20
- Run for the same amount of physical time (less epochs for red)



Results

- Network play style after 2000 epochs:
 - Connects stones (good)
 - Tries to capture (good)
 - Tries to fill the whole board (questionable)
 - Does not form Two Eyes (very very bad)
- Steady improvement each 100 epochs



- Example of white having two eyes. Black cannot capture the white group



Discussion


- Clearly, the bigger N , the better. Significant increase in performance with bigger N ?
 - Better preservation of good mutations
- Lower learning rate prevents divergence
- Play style could become proper over time.
- *"If you learn from idiots, you will remain an idiot"*
- Maybe better play style could be "forced" by changing reward function?



Discussion

- We tried to modify reward function.
- Reward faster play
 - Should result in less “board filling”
 - Resulted in networks giving up (passing) too early
- Reward bigger score
 - Resulted in more “infinite games”

- In an act of desperation, we tried to train from start.
- $N = 6$, learning_rate = 0.0001
- Reward function rewards fast play.
- Top player is selected and cloned 5 times.
- Result shows win rate against previously trained Stage 2 net over 60 epochs.





Playing example

Circle - Old, Cross - New


```
- 0 0 0 X X - X -  
0 0 0 X X X X X 0  
0 0 0 0 X X X X 0  
0 0 0 0 X X X - X  
0 0 0 0 0 X 0 0 X  
- 0 0 0 0 0 X X X  
0 0 0 0 X X X X X  
0 0 X X X X X X X |  
0 X X - X X X X X  
o is thinking [...]  
o selects: [0, 5]
```

```
- 0 0 0 X X - X -  
0 0 0 X X X X X 0  
0 0 0 0 X X X X 0  
0 0 0 0 X X X - X  
0 0 0 0 0 X 0 0 X  
0 0 0 0 0 0 X X X  
0 0 0 0 X X X X X  
0 0 X X X X X X X  
0 X X - X X X X X  
x is thinking [...]  
x selects: [0, 0]
```

```
X - - - X X - X -  
- - - X X X X X 0  
- - - - X X X X 0  
- - - - X X X - X  
- - - - - X 0 0 X  
- - - - - - X X X  
- - - - X X X X X  
- - X X X X X X X  
- X X - X X X X X
```

Discussion

- Correction, better play style CAN be forced, but reward function must be such from start.
- It's doable to train with small N but you have to "get lucky".
- Bigger N =>
 - More minima are explored and evaluated simultaneously.
 - Bigger probability of getting to global minima.
 - Better chance of preserving good mutations





Conclusion/Summary

More training needed

Bigger N (100-1000) should be used ?

-> Needs stronger hardware - a single GPU is not enough

Need deeper network?

Q&A

